

How would you deal with a customer complaint?

- Restore trust, fix root cause, prevent repeats
- Acknowledge quickly, get specifics: what they tried, expected, saw
- Reproduce, scope impact
- Tell them: what I know, what I do not know yet, next update time
- Fix with guardrails
- Debug, keep monitoring
- Follow up: root cause, what we changed
- Compye.ai: added Sentry to detect sooner and reduce repeat reports, then throttling and caching to reduce peak failures

Describe a complex problem step by step

- Columbia: co built machine driven benchmark for LLM chatbots in LGBTQ suicide prevention
- Stakes: supportive sounding response can still miss risk assessment, crisis resources, identity affirmation
- Personal: put myself in the client's shoes
- Built rigorous, scalable evaluation system
- Compared to clinician written reference
- Scored across multiple dimensions, not one metric
- Clinical alignment: ROUGE and METEOR vs human reference
- ROUGE: wording
- METEOR: detect paraphrases
- Result: single score is misleading

- Similar to GS AI Assistant rollout, emphasis on safe use inside the company

Why are you interested in Engineering at Goldman Sachs?

- Firas debbech vp and high pressure
- Engineering decisions matter, systems must be fast, reliable, safe
- Came to US as international student, later asylum, learned calm under pressure, goals where we cannot afford failures
- Goldman builds platforms teams and clients rely on
- Marquee: analytics through API for clients
- Legend: make data understandable, democratize it
- Same idea as my research pipeline
- Took “safe and affirming” and made defined dimensions, consistent inputs, repeatable scoring
- I will be effective because I bring measurable engineering
- Personally care about tech people can trust with finances, and in an abstract way, their life, mirrors my work so far

Situation where you had to analyze information to identify trends

- Research: evaluated multiple chatbots on LGBTQ suicide prevention prompts
- Single summary scores hid important failures
- Analyzed large evaluation output to find failure patterns
- Structured outputs to slice by metric and by prompt section
- Goldman parallel: Marquee turns large volumes of data into signals users can act on

- My signal was safety and inclusivity, not market data, engineering mindset identical
- Explained differences clearly and defensibly, not anecdotes
- Built repeatable way to re measure trends as models changed

Describe a problem you faced and how you solved it

- Internship: recurring user reported errors, performance spikes, poor visibility into what was breaking
- Task: reduce repeat failures, improve reliability, detect issues before complaints
- Added production error monitoring and alerts for visibility into exceptions and frequency
- Implemented request controls and caching to reduce repeated expensive calls, smooth spikes during high traffic
- Goldman parallel: reliability and predictable performance matter because users depend on systems, including platforms like Marquee
- Result: reduced time to detect issues, lowered repeat failures, stabilized peak performance

How would you deal with a customer complaint?

- Restore trust quickly, fix root cause safely, prevent repeats, communicate clearly the whole time
- Acknowledge quickly, get specifics
- Reproduce, scope impact
- Follow up: root cause, prevention steps

Describe a thematic project to produce high quality results

- Situation: Columbia research project to evaluate LLM chatbots for LGBTQ suicide prevention, where quality equals safety, trust, and clarity
- Task: produce rigorous, repeatable results that held up across models, not a one off demo
- Action: standardized the prompt format, compared against a clinician reference, and scored outputs across multiple dimensions, including ROUGE, METEOR, ethics, sentiment, inclusivity, and readability
- Result: high quality meant we could **explain strengths and failures with evidence**, and rerun the pipeline as models changed
- Goldman parallel: same mindset as Legend, define standards once so outputs are consistent, and GS AI Assistant, deploy GenAI with controls and evaluation

Describe a presentation you had to deliver and how you finished it

- Situation: I had to present our evaluation pipeline and results to a mixed audience, including people who were not deep in ML
- Task: make the work clear, credible, and actionable, without overwhelming them
- Action: I structured the talk around the six dimensions, used a simple story of why single scores mislead, then walked through a table of model tradeoffs and what we would do next
- Result: the audience left with clear takeaways, what was strong, what was risky, and what to improve
- Goldman parallel: similar to explaining Marquee analytics, the goal is turning complexity into decisions people can act on

Classmate depending on me to finish a project

- Situation: in a team setting, others were blocked on getting results finalized and reproducible

- Task: take ownership of the critical path so the group could submit on time
- Action: I locked the scope, broke work into small deliverables, finished the evaluation script and results output, and documented how to rerun everything so no one depended on my laptop
- Result: we shipped on time and the work was reproducible for everyone
- Goldman parallel: similar to platform work like GS Quant and Legend, the win is making outputs reusable by the whole team, not just the person who built it

Describe a mistake and a failure and what you learned from it

- Situation: early in the project, I relied too much on automated scoring and assumed a metric captured what we meant
- Task: correct the failure and prevent misleading conclusions
- Action: I rechecked edge cases, found where the metric was blind, tightened the definitions, and added checks so results could not look “good” while missing important signals
- Result: we improved the integrity of the evaluation and reduced the risk of false confidence
- Lesson: metrics are tools, not truth, and high stakes systems need validation and clear assumptions
- Goldman parallel: same risk mindset as finance analytics, a clean number is useless if it hides model risk

Strongest skill and how you developed it

- Strongest skill: staying calm and executing step by step under pressure, while keeping work measurable
- How I developed it: my background forced me to be resilient, then engineering trained me to turn uncertainty into plans, metrics, and iteration

- Proof: I build evaluation harnesses, monitoring loops, and clear definitions so decisions are based on evidence
- Goldman parallel: this fits engineering at Goldman, where you need reliability and speed without losing correctness

Provide a recent example of a new skill and how you started

- Skill: building machine driven evaluation for LLM behavior, including sentiment alignment and automated scoring
- How I started: I read prior benchmarks, defined what we needed, built a first version quickly, then improved it through repeated testing and error analysis
- Result: I could evaluate many models consistently and explain tradeoffs clearly
- Goldman parallel: same learning loop you need for GS AI Assistant work, ship safely, measure, then iterate

If I don't have the experience for a task, what do I do

- I clarify the goal, success metrics, and constraints first
- I break it into small pieces and start with a minimal version that proves the path
- I ask for a quick check from someone experienced, early, not after I build the wrong thing
- I read the docs and existing codebase patterns, then follow them
- I communicate progress and risks, and I keep tightening until it meets the bar
- Goldman parallel: this is how you ramp on large platforms like Legend or Marquee, learn the standards, ship something small, then scale it safely

I used Hugging Face to access LLMs like ChatGPT and Claude, then called them through Python to generate clinical style responses for LGBTQ suicide prevention prompts. I wrote a

PyTorch based evaluation program that scored each response with metrics like ROUGE and METEOR, and only returned the highest scoring response based on the benchmark results.\

Redis for the live exam timer and PostgreSQL for storing attempts; I used an LLM API to generate grading feedback and explanations, and I used Sentry to catch errors and latency issues in real time.

Docker packages your app and everything it needs into a container so it runs the same way on your laptop

Django is a Python web framework

I added Sentry, which automatically reports crashes and slowdowns, so I could see problems right away

I added a speed limit and a memory layer using Redis so the system did not call the LLM too many times

I wrote a small Python wrapper around the ChatGPT API that took a prompt, called the model, and returned the response text I passed each response into my benchmark functions, like ROUGE, METEOR saved the scores for comparison and then give best answer

Marquee

Marquee is Goldman's online platform for big professional clients. It puts market data, analysis, and tools in one place, and it can connect to a client's systems through APIs.

GS AI Assistant

GS AI Assistant is Goldman's internal ChatGPT style helper for employees. It helps with things like summarizing, drafting, and coding, and it is built to handle sensitive work safely.

Legend

Legend is Goldman's system for organizing and standardizing data. It makes sure teams use the same definitions, so data stays consistent across the firm.

Legend AI Query

Legend AI Query lets employees ask for data in plain English. It helps you find the right data without knowing exact table names.

Legend Copilot

Legend Copilot is an AI helper for people building data models in Legend. It speeds up creating and updating data definitions and related APIs.